

CS4495/6495



Introduction to Computer Vision

4C-L1 *Robust error functions*

Feature-based alignment to find transforms

Overall strategy:






1. Compute features – detect and describe



Feature-based alignment to find transforms

Overall strategy:







1. Compute features – detect and describe



Feature-based alignment to find transforms

Overall strategy:


2. Find some useful matches:
Kd-tree, Best-Bin, Hashing



Feature-based alignment to find transforms

Overall strategy:

3. **Compute** and apply the best transformation:
e.g. affine, translation, or homography




Feature-based alignment to find transforms

Overall strategy:

3. Compute and **apply** the best transformation:
e.g. affine, translation, or homography

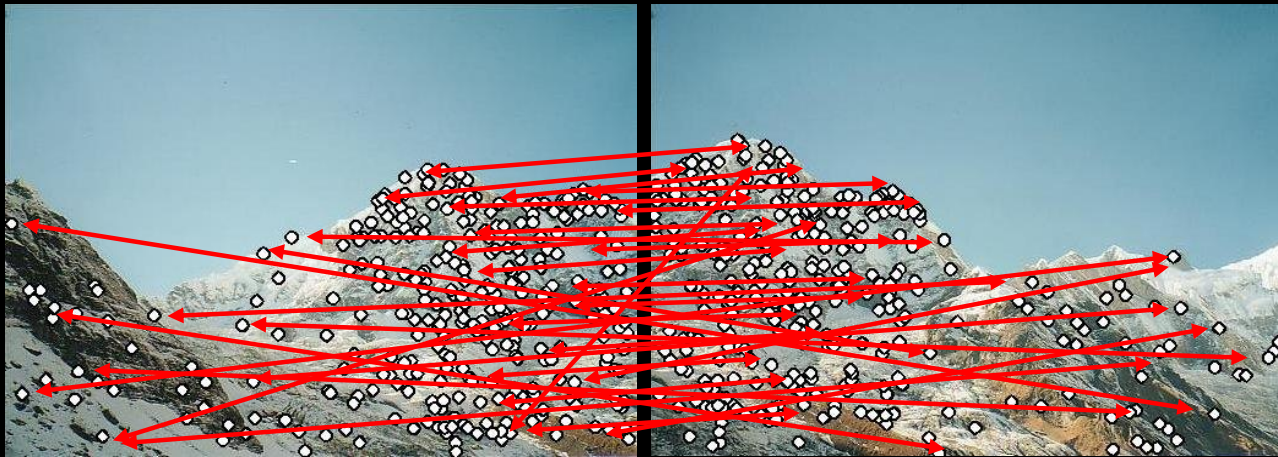


Feature-based alignment algorithm



1. Extract features


Feature-based alignment algorithm



2. Compute *putative matches* – e.g. “closest descriptor”

Kd-tree, best bin, etc...


Feature-based alignment algorithm



3. Loop until happy:

- *Hypothesize* transformation T from some matches
- *Verify* transformation (search for other matches consistent with T) – mark best

Feature-based alignment algorithm



4. Apply best transformation.

How to get “putative” matches?

Feature matching

- Exhaustive search
- Hashing
- Nearest neighbor techniques


.... but these give the best match. How do we know it's a good one?

Feature-space outlier rejection

- Let's not match all features, but only these that have “similar enough” matches?
- How can we do it?
 - $\text{SSD}(\text{patch1}, \text{patch2}) < \text{threshold}$
 - How to set threshold?


Feature-space outlier rejection

- How to set threshold?



A better way [Lowe, 1999]:

- 1-NN: SSD of the closest match
- 2-NN: SSD of the second-closest match
- Look at how much better 1-NN is than 2-NN, e.g. 1-NN/2-NN




Feature matching

- Exhaustive search
- Hashing
- Nearest neighbor techniques
- But...remember the distinctive vs invariant competition? Implies:
 - *Problem: Even when pick best match, still lots (and lots) of wrong matches – “outliers”. What should we do?*

Model Fitting

- Choose a parametric model to represent a set of features – *remember this???*



simple model: lines



simple model: circles

Fitting: Issues

Case study: Line detection


- **Noise** in the measured feature locations
- **Extraneous data**: clutter (outliers), multiple lines
- **Missing data**: occlusions



Typical least squares line fitting


- Data: $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation: $y_i = m x_i + b$
- Find (m, b) to minimize:

$$E = \sum_{i=1}^n (y_i - m x_i - b)^2$$



Typical least squares line fitting

$$E = \sum_{i=1}^n (y_i - m x_i - b)^2$$



$$E = \sum_{i=1}^n \left(y_i - \begin{bmatrix} x_i & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right)^2 = \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right\|^2 = \|\mathbf{y} - \mathbf{X}\mathbf{h}\|^2$$

\mathbf{h}

$$E = (\mathbf{y} - \mathbf{X}\mathbf{h})^T (\mathbf{y} - \mathbf{X}\mathbf{h}) = \mathbf{y}^T \mathbf{y} - 2(\mathbf{X}\mathbf{h})^T \mathbf{y} + (\mathbf{X}\mathbf{h})^T (\mathbf{X}\mathbf{h})$$


Typical least squares line fitting

$$E = (\mathbf{y} - \mathbf{X}\mathbf{h})^T (\mathbf{y} - \mathbf{X}\mathbf{h}) = \mathbf{y}^T \mathbf{y} - 2(\mathbf{X}\mathbf{h})^T \mathbf{y} + (\mathbf{X}\mathbf{h})^T (\mathbf{X}\mathbf{h})$$

$$\Rightarrow \frac{dE}{d\mathbf{h}} = 2\mathbf{X}^T \mathbf{X}\mathbf{h} - 2\mathbf{X}^T \mathbf{y} = 0$$

$$\mathbf{X}^T \mathbf{X}\mathbf{h} = \mathbf{X}^T \mathbf{y} \Rightarrow \mathbf{h} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$\underbrace{(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T}_{\text{pseudoinverse}} \mathbf{y}$




*Standard over-constrained
least squares solution*

Problem with “vertical” least squares

- Not rotation-invariant
- Fails completely for vertical lines

Total least squares


- Distance between point (x_i, y_i) and line $ax + by = d$
- Find (a, b, d) to minimize the sum of squared *perpendicular* distances



$$E = \sum_{i=1}^n (a x_i + b y_i - d)^2$$

Total least squares

$$E = \sum_{i=1}^n (a x_i + b y_i - d)^2$$




$$\frac{\partial E}{\partial d} = \sum_{i=1}^n -2(a x_i + b y_i - d) = 0$$

$$\Rightarrow d = \frac{a}{n} \sum_{i=1}^n x_i + \frac{b}{n} \sum_{i=1}^n y_i = a \bar{x} + b \bar{y}$$

Total least squares

$$E = \sum_{i=1}^n (a x_i + b y_i - d)^2$$

$$d = a \bar{x} + b \bar{y}$$



$$E = \sum_{i=1}^n (a (x_i - \bar{x}) + b (y_i - \bar{y}))^2 = \left\| \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} \right\|^2 = (\mathbf{U} \mathbf{h})^T (\mathbf{U} \mathbf{h})$$

\mathbf{h}

$$\frac{dE}{d \mathbf{h}} = 2(\mathbf{U}^T \mathbf{U}) \mathbf{h} = 0$$

Total least squares


$$E = \sum_{i=1}^n (a x_i + b y_i - d)^2$$

$$d = a \bar{x} + b \bar{y}$$

$$\frac{dE}{d\mathbf{h}} = 2(U^T U)\mathbf{h} = 0$$


Solution to $(U^T U)\mathbf{h} = 0$, subject to $\|\mathbf{h}\|^2 = 1$:

eigenvector of $U^T U$ associated with the smallest eigenvalue
(Again SVD to least squares solution to *homogeneous linear system*)



Least squares as likelihood maximization

- **Generative model:** Line points are corrupted by Gaussian noise in the direction perpendicular to the line




Least squares as likelihood maximization

$$\begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} u \\ v \end{pmatrix} + \varepsilon \begin{pmatrix} a \\ b \end{pmatrix}$$

point on
the line


noise:
sampled from
zero-mean
Gaussian with
std. dev. σ

normal
direction



Least squares: Non-robustness to (very) non-Gaussian noise


- Least squares fit to the red points:



Least squares: Non-robustness to (very) non-Gaussian noise

- Least squares fit with an outlier...

Problem: *squared error heavily penalizes outliers*




Robust estimators

General approach: minimize $\sum_i \rho \left(r_i \left(x_i, \theta \right); \sigma \right)$

$r_i(x_i, \theta)$ – residual of i^{th} point w.r.t.
model parameters θ


ρ – robust function with scale
parameter σ

Robust estimators




The robust function ρ behaves like squared distance for small values of the residual u but saturates for larger values of u

Choosing the scale: Just right




The effect of the outlier is minimized

Choosing the scale: Too small



Choosing the scale: Too large



Behaves much the same as least squares